

ECE 3574: Thread Synchronization

Changwoo Min

Thread Synchronization

- Today we are going to look at how to manage access to shared memory using a mutex and how to build higher-level abstractions, a semaphore and a thread-safe queue.
 - Races and Atomics
 - Mutex's and locking
 - Condition Variables
 - Semaphore
 - Building a Semaphore using C++11
 - QSemaphore

Recall the `QSharedMemory` class had a lock/unlock mechanism

- Why was this needed?
- How is such a thing implemented?

The Why Question: Data races

```
// Consider two threads that share an integer pointer, x with a loop  
// (see race.cpp)
```

```
while(*x > 0){  
    std::this_thread::sleep_for(std::chrono::nanoseconds(100));  
    *x -=1;  
}
```

```
// where the sleep is a stand-in for some computational work.  
// - what is the value of *x after the threads execute?  
// We say the threads are racing to get to the value of *x == 0.
```

Synchronization is built on the idea of atomics

- An atomic operation is one that is guaranteed to not cause data races.
- Example `int` vs `atomic_int`

```
int x;  
x = anothervar;
```

// v/s

```
std::atomic_int x;  
x.store(anothervar);
```

*// The latter is compiled to instructions that lock the memory bus
// during the assignment.
// How this works at the hardware level is complicated due to cache
// lines etc.*

Atoms can be used directly or form the basic of a locking mechanism

- create an atomic boolean initialized to false
- to lock test if the value is false and if so set it to true (exchange), else try again (AKA test and set). access the locked resource
- to unlock, set the atomic bool back to false
- See `race_atomic.cpp`
 - [`std::atomic`](#), [`ATOMIC_FLAG_INIT`](#)
- Lucky for us, more high-level locking semantics are defined in the threading library

A mutex (MUTual EXclusion) is an object with exclusive ownership semantics

- It provides a synchronization primitive for protecting shared memory from simultaneous writes and/or reads.
 - In the case of IPC the memory being protected is shared memory between processes.
 - In the case of threads the memory being protected is the heap.
- I will use threads to describe the ideas, but this works for IPC shared memory as well.

A mutex has two states: locked and unlocked

- Multiple threads may share a mutex variable, but only one can own it at a time. (mutual exclusion)
- To gain ownership a thread locks the mutex. If already locked by another thread this blocks. (lock)
- To release ownership a thread unlocks the mutex. (unlock)
- Typically a thread can also try to lock a mutex getting a bool flag indicating success/failure. This does not block. (try lock)

Basic protection of object using a mutex

0. Associate a mutex with the object.
1. Before accessing the object, lock the mutex.
2. Perform the access (read or write).
3. Unlock the mutex

All access goes through this lock/unlock sequence. If you forget to unlock you get a deadlock, and that object cannot be accessed.

Deadlocks

- Situations in which one or several threads are waiting on locks for one or several resources that will never be freed
 - None of the threads can continue
- Self-deadlock
 - non-recursive lock

```
acquire lock
acquire lock, again
wait for lock to become available
...
```

Deadlocks

- **Deadly embrace (ABBA deadlock)**

Thread 1	Thread 2
acquire lock A	acquire lock B
try to acquire lock B	try to acquire lock A
wait for lock B	wait for lock A

- **Deadlock prevention: lock ordering**
 - Nested locks must *always* be obtained in the *same order*.
 - This prevents the deadly embrace deadlock.

`std::mutex` in C++11

- `lock()` : locks the mutex, blocks if the mutex is not available
- `try_lock()` : try to lock the mutex, returns false if the mutex is not available
- `unlock()` : unlocks the mutex
- Failing to unlock causes a deadlock.
- See `simple_mutex_ex.cpp`.
- Reference: [link](#)

`std::lock_guard` in C++11

- The mutex is a resource that requires careful handling (e.g. to prevent deadlocks). The [C++ RAII](#) mechanism is ideal for this.
 - lock in a constructor
 - unlock in a destructor
 - let stack allocation handle the duration of the lock
- This can prevent many deadlocks, particularly those caused by an exception interrupting the lock process.
- This is what [std::lock_guard](#) does. It cannot be locked/unlocked outside its constructor/destructor. See `lock_guard_ex.cpp`.

`std::unique_lock` in C++11

- `std::unique_lock` is a more sophisticated wrapper around a `std::mutex`. Adds:
 - RAI lock/unlock like `lock_guard`
 - [deferred locking \(for simultaneous locking of multiple mutex's\)](#)
 - time-constrained `try_lock`: `try_lock_for` and `try_lock_until`
 - recursive locking
 - transfer of lock ownership
 - ability to use with condition variables
- See `unique_lock_ex.cpp`.

Condition variable

- A condition variable, and its associated mutex, allows multiple threads to communicate by one thread notifying others they can proceed
 - [std::condition_variable](#)

Condition variable

- Suppose multiple threads are sharing a variable with an associated `std::mutex`. A thread that wants to access the variable:
 - locks the mutex
 - reads or updates the shared variable
 - unlocks the mutex
 - calls `notify_one` or `notify_all` method of the `std::condition_variable` object

Condition variable

- A thread waiting on the notification (via a `std::condition_variable`):
 - instantiates a `unique_lock` on the shared variable's mutex, but does not try to lock it
 - instead it calls `wait`, `wait_for`, or `wait_until` method, suspending the thread (possibly with a timeout)

The condition variable receives a notification

- when
 - another thread calls notify
 - a timeout expires
 - a spurious wakeup occurs
- Upon notification the thread is awakened, and the mutex acquired. You should check the condition and call wait again in case the notification was spurious.

The condition variable receives a notification

- Spurious wakeups are a bit mysterious. Why would the notification be sent if the condition was not true?
 - you might have a bug in the other thread
 - there are also performance-related reasons why checking in the thread implementation is not done.
- See `condition_variable_ex.cpp`.

Semaphores

- A **semaphore** is an abstraction of an integer that can be used to share resources between threads.
- A threshold (default of 0) can be used to allow multiple threads access, but limit it.
 - use `Semaphore::up` to release a resource, increments the integer
 - use `Semaphore::down` to acquire a resource, decrements the integer (blocks until above threshold).

Semaphores

- `up` is also called `release`, `down` is also called `acquire`.
- A semaphore can be used to give threshold number of worker threads access to a resource at a time. A common example is limited file-system IO in distributed systems.
- See `test_semaphore.cpp`.

Building a semaphore using C++11

- See semaphore.h and semaphore.cpp.

Qt has a built in QSemaphore

- The constructor creates a semaphore guarding n resource units (by default, 0).
- `acquire(int n = 1)`, acquires n resource units, blocking until n are available
- `available()` returns the number of resources available
- `release(int n = 1)` releases n resource units
- There are try versions of acquire that return immediately on failure or after a timeout.

Milestone 3: FAQ

- Drawing a table in Qt
 - [QStandardItemModel](#)
 - [QTableView](#)
 - [Examples](<http://doc.qt.io/qt-5/model-view-programming.html>)
- Highlighting a text
 - [QTextCursor](#)
 - [QTextCharFormat](#)

Next Actions and Reminders

- Milestone 3 due Monday 4/9 by 11:59 pm.